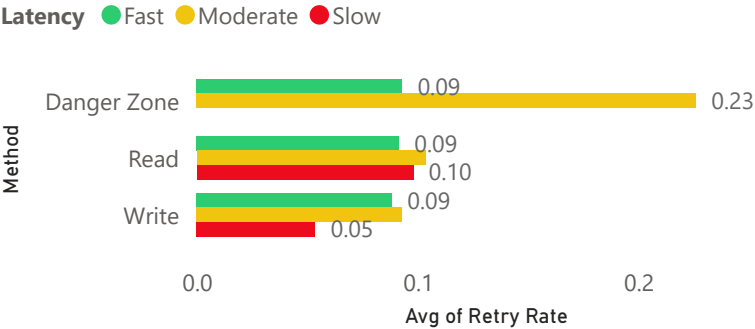


⏪ System Performance by Latency Range & Retry Behaviour

Latency

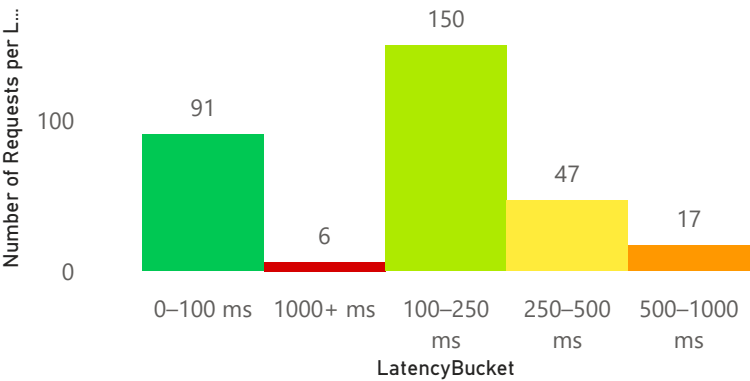
All

Which Operations Retry Most—and When Performance Drops



Retry Behaviour by Method and Latency: Highlighting which operations trigger client retries across performance tiers.

Number of Requests per Latency Range by LatencyBucket



- Fast (0-250 ms)
- Moderate (250-500 ms)
- Slower (500-1000 ms)
- High Latency (1000+ ms)

Requests	Per Second	RSP (ms)	p95 (ms)	% Error Rate	Avg Retry Rate
5K	0.06	153	233.00	14.18	0.09

Infrastructure Health: Server Load & Failures

Visualizing server saturation via CPU/memory usage, traffic volume, and failure rates

Server ID	Avg Memory	Avg RPS	Total Failures	Avg CPU
server-003	52.56	69	1	58.55
server-005	36.84	48	1	42.01
server-008	53.05	70	2	55.65
server-013	49.97	67	2	56.61
server-016	45.78	60	2	53.64
server-017	45.47	57	2	52.25
server-020	46.27	59	2	58.92
server-002	42.60	59	3	53.94
server-004	47.80	69	3	56.38
server-009	51.27	67	3	60.30
server-011	41.08	62	3	55.77
server-018	53.68	69	3	55.27
server-007	43.32	60	4	48.38
server-010	46.55	72	4	59.47
server-015	57.30	85	4	62.38
server-001	48.31	67	5	61.17
server-012	38.78	59	5	43.03
server-006	45.56	73	6	60.14
server-014	40.07	65	9	49.93
Total	932.07	1288	64	1094.90

Server Load & Failure Overview: Mapping average CPU, memory, and total failures per node to assess system pressure.


High failure counts and CPU usage suggest overburdened or unstable instances. Use this panel to prioritize scaling or mitigation.

Region

All


Reliability vs Performance: Mapping anomaly severity, retry load, and failure type across operational categories.

Larger bubbles = higher retry pressure. Rightward = more anomalies. Higher = slower response



Anomaly Severity Timeline (Per Minute)

Grouped by Latency & Method Category



Severity Trends by Minute: Visualizing anomaly density across request time, segmented by risk level and traffic behaviour

Higher scores indicate more severe anomalies. Coloured lines show tiered risk levels based on latency, method type, and error flags.

1.00

%

153.10

Latency

0.09

Retry Rate